

```
In [225]: 1 import sqlalchemy
2 from sqlalchemy import create_engine
3 import numpy as np
4 import pandas as pd
5 import statsmodels.api as sm
6 import matplotlib.pyplot as plt
7 %config InlineBackend.figure_format = "svg"
8 from sklearn.model_selection import train_test_split
9 from sklearn.linear_model import LogisticRegression
10 from sklearn.metrics import accuracy_score, classification_report
11 import seaborn as sns
12 sns.set(style="white")
13 sns.set(style="whitegrid", color_codes=True)
```

Charlie Morris - Cosc 61 Final Project

```
In [226]: 1 #Connect to MySQL database called 'baseball'
2 con_string = 'mysql+pymysql://root:C&bDanHal@127.0.0.1:3306/baseball'
3 engine = create_engine(con_string)
```

```
In [227]: 1 query2 = """
2         SHOW TABLES
3         """
4
5 df2_read_sql = pd.read_sql(query2, engine)
6 df2_read_sql
```

Out[227]:

Tables_in_baseball	
0	award_types
1	awards
2	batting_stats
3	countries
4	errors
5	franchises
6	hall_of_fame
7	pitching_stats
8	players
9	positions
10	salaries
11	team_stats
12	teams
13	world_series

Logistic Regression: Probability of Making HOF Vesus Career Pitching SOs

```
In [228]: 1 #Grab the data for pitchers and their career strikeouts
          2 #Look only at pitchers with minimum 201 strikeouts
          3
          4 query = """
          5     SELECT playerID, SUM(SO) AS Career_SO
          6     FROM pitching_stats
          7     WHERE playerID IN
          8         (SELECT playerID
          9         FROM players
10         WHERE '1900-12-31' < finalGame AND finalGame < '2008-12-31')
11     GROUP BY playerID
12     HAVING Career_SO > 200;
13 """
          14
          15 df_CareerSOs = pd.read_sql(query, engine)
          16 df_CareerSOs
```

Out[228]:

	playerID	Career_SO
0	aasedo01	641.0
1	abbotgl01	484.0
2	abbotji01	888.0
3	abbotpa01	496.0
4	abernte02	765.0
...
1935	zambrvi01	529.0
1936	zimmeje02	213.0
1937	zoldasa01	207.0
1938	zuberbi01	383.0
1939	zuverge01	223.0

1940 rows × 2 columns

```
In [229]: 1 #Grab data on HOF players
          2 query = """
          3     SELECT *
          4     FROM hall_of_fame;
          5     """
          6
          7 df_HOF = pd.read_sql(query, engine)
          8 df_HOF
```

Out[229]:

	playerID	year
0	aaronha01	1982
1	alexape01	1938
2	alomaro01	2011
3	alstowa01	1983
4	andersp01	2000
...
282	wynnea01	1972
283	yastrca01	1989
284	youngcy01	1937
285	youngro01	1972
286	yountro01	1999

287 rows × 2 columns

```
In [230]: 1 #Merge the 2 data frames
          2 #1 means made HOF, 0 means didn't make HOF
          3 full_df = pd.merge(df_CareerSOs, df_HOF, on = "playerID", how = "left")
          4 full_df['HOF_Status'] = full_df['year'].notnull().astype(int)
          5 full_df
```

Out[230]:

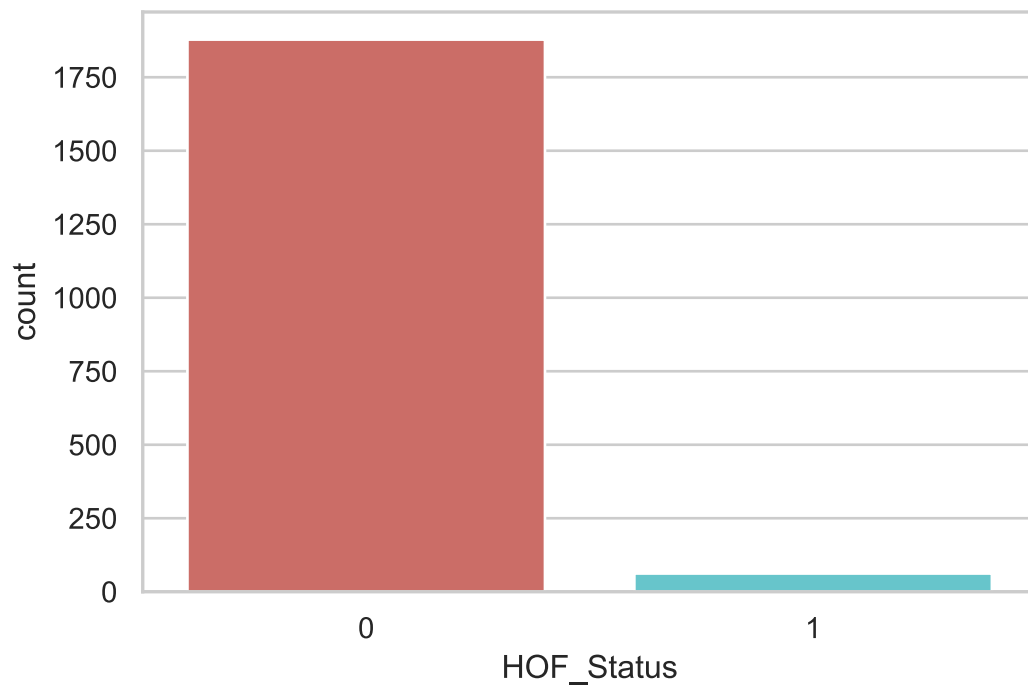
	playerID	Career_SO	year	HOF_Status
0	aasedo01	641.0	NaN	0
1	abbotgl01	484.0	NaN	0
2	abbotji01	888.0	NaN	0
3	abbotpa01	496.0	NaN	0
4	abernte02	765.0	NaN	0
...
1935	zambrvi01	529.0	NaN	0
1936	zimmeje02	213.0	NaN	0
1937	zoldasa01	207.0	NaN	0
1938	zuberbi01	383.0	NaN	0
1939	zuverge01	223.0	NaN	0

1940 rows × 4 columns


```
In [231]: 1 #Count how many pitchers in the HOF
          2 full_df['HOF_Status'].value_counts()
```

```
Out[231]: 0    1878
          1     62
          Name: HOF_Status, dtype: int64
```

```
In [232]: 1 #Plot HOF players count and non-HOF players count
2 sns.countplot(x='HOF_Status', data = full_df, palette='hls')
3 plt.show()
```



```
In [233]: 1 #Scatterplot of data
2 sns.scatterplot(x="Career_SO", y = "HOF_Status", data = full_df)
3 plt.show()
```



```
In [234]: 1 #Mean values of each group
          2 full_df.groupby('HOF_Status').mean()['Career_SO']
```

```
Out[234]: HOF_Status
0      630.900958
1     2090.032258
Name: Career_SO, dtype: float64
```

```
In [235]: 1 #Run the logistic regression
          2 X = full_df[['Career_SO']]
          3 y = full_df['HOF_Status']
          4
          5 # Split the data into training and testing sets
          6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
          7
          8 # Initialize the logistic regression model
          9 model = LogisticRegression()
         10
         11 # Train the model on the training data
         12 model.fit(X_train, y_train)
         13
         14 # Make predictions on the test data
         15 y_pred = model.predict(X_test)
         16
         17 # Evaluate the model's accuracy
         18 accuracy = accuracy_score(y_test, y_pred)
         19 print(f"Accuracy: {accuracy:.2f}")
         20
         21 # Print classification report
         22 print(classification_report(y_test, y_pred))
```

Accuracy: 0.98

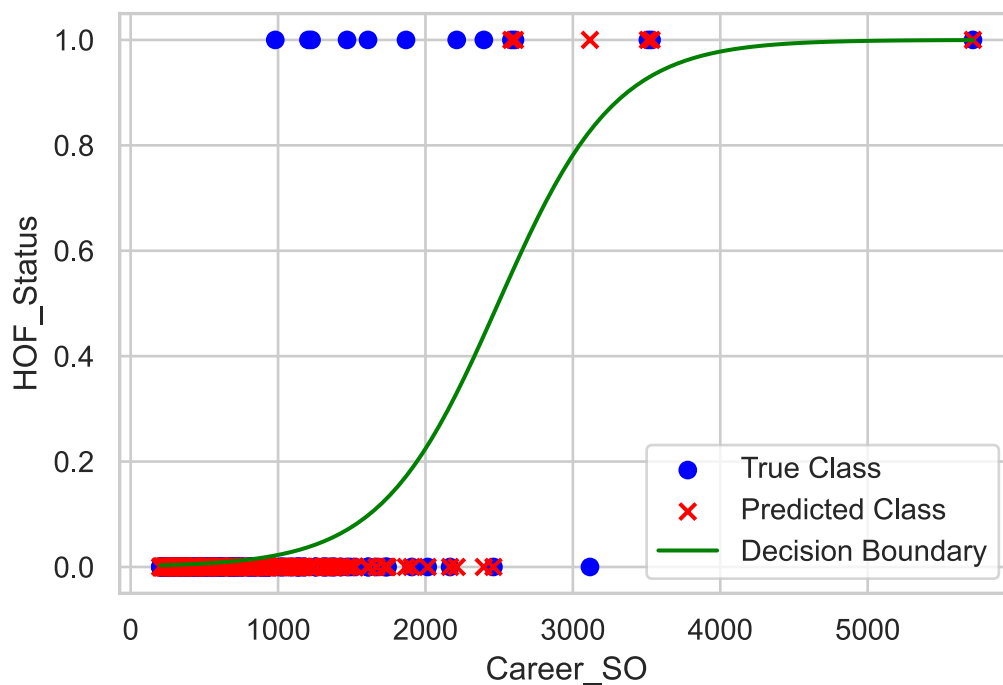
	precision	recall	f1-score	support
0	0.98	1.00	0.99	375
1	0.83	0.38	0.53	13
accuracy			0.98	388
macro avg	0.91	0.69	0.76	388
weighted avg	0.97	0.98	0.97	388

```

In [236]: 1 # Plot the data points and decision boundary
2 plt.scatter(X_test, y_test, color='blue', label='True Class')
3 plt.scatter(X_test, y_pred, color='red', marker='x', label='Predicted C
4 plt.xlabel('Career_SO')
5 plt.ylabel('HOF_Status')
6 plt.legend()
7
8 # Create a range of values for the x-axis
9 x_range = np.linspace(X.min(), X.max(), num=100)
10 # Calculate the corresponding y-values using the logistic regression mo
11 y_range = model.predict_proba(x_range.reshape(-1, 1))[:, 1]
12
13 # Plot the decision boundary
14 plt.plot(x_range, y_range, color='green', label='Decision Boundary')
15 plt.legend()
16
17 plt.show()

```

/Users/charliemorris/opt/anaconda3/lib/python3.9/site-packages/sklearn/base.py:450: UserWarning: X does not have valid feature names, but Logistic Regression was fitted with feature names
 warnings.warn(



Logistic Regression: Probability of Making HOF Vesus Career Batting Average

```
In [237]: 1 #Grab the data for batters and their career batting averages
2 #Look only at batters with minimum 1500 at bats
3
4 query = """
5     SELECT playerID, ROUND(SUM(H)/SUM(AB) * 1000) AS Career_BA
6     FROM batting_stats
7     WHERE playerID IN
8         (SELECT playerID
9          FROM players
10         WHERE '1930-12-31' < finalGame AND finalGame < '2008-12-31')
11     GROUP BY playerID
12     HAVING SUM(AB) > 2000;
13 """
14
15 df_CareerBA = pd.read_sql(query, engine)
16 df_CareerBA
```

Out[237]:

	playerID	Career_BA
0	aaronha01	305.0
1	abbotku01	256.0
2	adairje01	254.0
3	adamsbo03	269.0
4	adamsbu01	266.0
...
1448	zarilal01	276.0
1449	zeileto01	265.0
1450	zernigu01	265.0
1451	zimmedo01	235.0
1452	ziskri01	287.0

1453 rows × 2 columns


```
In [238]: 1 #Grab data on HOF players
          2 query = """
          3     SELECT *
          4     FROM hall_of_fame;
          5     """
          6
          7 df_HOF = pd.read_sql(query, engine)
          8 df_HOF
```

Out[238]:

	playerID	year
0	aaronha01	1982
1	alexape01	1938
2	alomaro01	2011
3	alstowa01	1983
4	andersp01	2000
...
282	wynnea01	1972
283	yastrca01	1989
284	youngcy01	1937
285	youngro01	1972
286	yountro01	1999

287 rows × 2 columns

```
In [239]: 1 #Merge the 2 data frames
          2 #1 means made HOF, 0 means didn't make HOF
          3 full_df = pd.merge(df_CareerBA, df_HOF, on = "playerID", how = "left")
          4 full_df['HOF_Status'] = full_df['year'].notnull().astype(int)
          5 full_df
```

Out[239]:


	playerID	Career_BA	year	HOF_Status
0	aaronha01	305.0	1982.0	1
1	abbotku01	256.0	NaN	0
2	adairje01	254.0	NaN	0
3	adamsbo03	269.0	NaN	0
4	adamsbu01	266.0	NaN	0
...
1448	zarilal01	276.0	NaN	0
1449	zeileto01	265.0	NaN	0
1450	zernigu01	265.0	NaN	0
1451	zimmedo01	235.0	NaN	0
1452	ziskri01	287.0	NaN	0

1453 rows × 4 columns


```
In [240]: 1 #Count how many batters in the HOF
          2 full_df['HOF_Status'].value_counts()
```

Out[240]: 0 1340
1 113
Name: HOF_Status, dtype: int64

```
In [241]: 1 #Plot HOF players count and non-HOF players count
2 sns.countplot(x='HOF_Status', data = full_df, palette='hls')
3 plt.show()
```



```
In [242]: 1 #Scatterplot of data
2 sns.scatterplot(x="Career_BA", y = "HOF_Status", data = full_df)
3 plt.show()
```



```
In [243]: 1 #Mean values of each group
          2 full_df.groupby('HOF_Status').mean()['Career_BA']
```

```
Out[243]: HOF_Status
0      266.487313
1      298.415929
Name: Career_BA, dtype: float64
```

```
In [244]: 1 #Run the logistic regression
          2 X = full_df[['Career_BA']]
          3 y = full_df['HOF_Status']
          4
          5 # Split the data into training and testing sets
          6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
          7
          8 # Initialize the logistic regression model
          9 model = LogisticRegression()
         10
         11 # Train the model on the training data
         12 model.fit(X_train, y_train)
         13
         14 # Make predictions on the test data
         15 y_pred = model.predict(X_test)
         16
         17 # Evaluate the model's accuracy
         18 accuracy = accuracy_score(y_test, y_pred)
         19 print(f"Accuracy: {accuracy:.2f}")
         20
         21 # Print classification report
         22 print(classification_report(y_test, y_pred))
```

Accuracy: 0.94

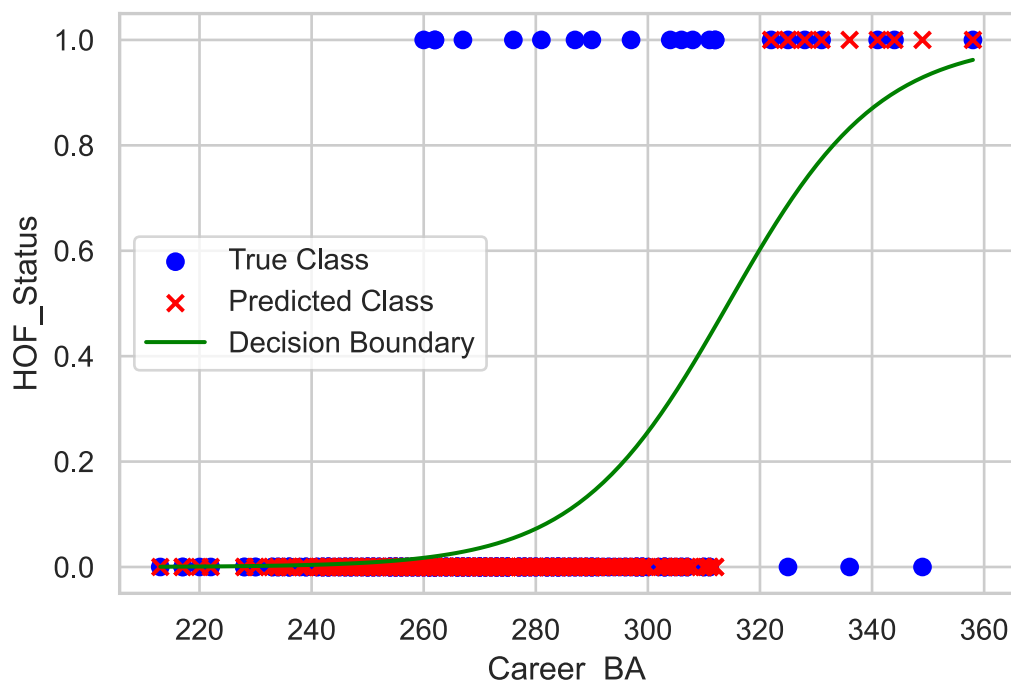
	precision	recall	f1-score	support
0	0.95	0.99	0.97	269
1	0.70	0.32	0.44	22
accuracy			0.94	291
macro avg	0.82	0.65	0.70	291
weighted avg	0.93	0.94	0.93	291

```

In [245]: 1 # Plot the data points and decision boundary
2 plt.scatter(X_test, y_test, color='blue', label='True Class')
3 plt.scatter(X_test, y_pred, color='red', marker='x', label='Predicted C
4 plt.xlabel('Career_BA')
5 plt.ylabel('HOF_Status')
6 plt.legend()
7
8 # Create a range of values for the x-axis
9 x_range = np.linspace(X.min(), X.max(), num=100)
10 # Calculate the corresponding y-values using the logistic regression mo
11 y_range = model.predict_proba(x_range.reshape(-1, 1))[:, 1]
12
13 # Plot the decision boundary
14 plt.plot(x_range, y_range, color='green', label='Decision Boundary')
15 plt.legend()
16
17 plt.show()

```

/Users/charliemorris/opt/anaconda3/lib/python3.9/site-packages/sklearn/base.py:450: UserWarning: X does not have valid feature names, but Logistic Regression was fitted with feature names
 warnings.warn(



Prime Age for Baseball

```
In [246]: 1 #Look at the baseball players with career HR totals of at least 275
2 #Each row is their age during a season and their HR totals
3 query = """
4     SELECT year - birthYear AS Age, HR
5     FROM players AS p
6     INNER JOIN batting_stats AS bs
7     ON p.playerID = bs.playerID
8     WHERE p.playerID IN
9         (SELECT playerID
10          FROM batting_stats
11          GROUP BY playerID
12          HAVING SUM(HR) > 275);
13 """
14
15 df = pd.read_sql(query, engine)
16 df
```

Out[246]:

	Age	HR
0	20	13
1	21	27
2	22	26
3	23	44
4	24	30
...
3078	32	18
3079	33	17
3080	34	6
3081	34	15
3082	35	0

3083 rows × 2 columns

```
In [247]: 1 counts = df['Age'].value_counts().sort_index()
          2 counts_df = counts.reset_index()
          3 counts_df.columns = ['Age', 'Count']
          4 counts_df
```

Out[247]:

	Age	Count
0	17	1
1	18	3
2	19	12
3	20	28
4	21	72
5	22	108
6	23	143
7	24	174
8	25	167
9	26	166
10	27	172
11	28	174
12	29	172
13	30	174
14	31	173
15	32	174
16	33	165
17	34	178
18	35	175
19	36	160
20	37	133
21	38	117
22	39	82
23	40	69
24	41	43
25	42	29
26	43	10
27	44	6
28	45	2
29	46	1

```
In [248]: 1 medianHRs = df.groupby('Age').median()["HR"]  
2 medianHRs_df = medianHRs.reset_index()  
3 medianHRs_df.columns = ['Age', 'MedianHRs']  
4 medianHRs_df
```

Out[248]:

	Age	MedianHRs
0	17	0.0
1	18	0.0
2	19	0.5
3	20	4.5
4	21	3.0
5	22	10.0
6	23	16.0
7	24	18.0
8	25	23.0
9	26	26.0
10	27	27.5
11	28	28.5
12	29	29.0
13	30	26.0
14	31	28.0
15	32	26.0
16	33	23.0
17	34	21.0
18	35	19.0
19	36	18.0
20	37	16.0
21	38	11.0
22	39	14.5
23	40	12.0
24	41	9.0
25	42	4.0
26	43	9.0
27	44	3.5
28	45	2.5
29	46	1.0

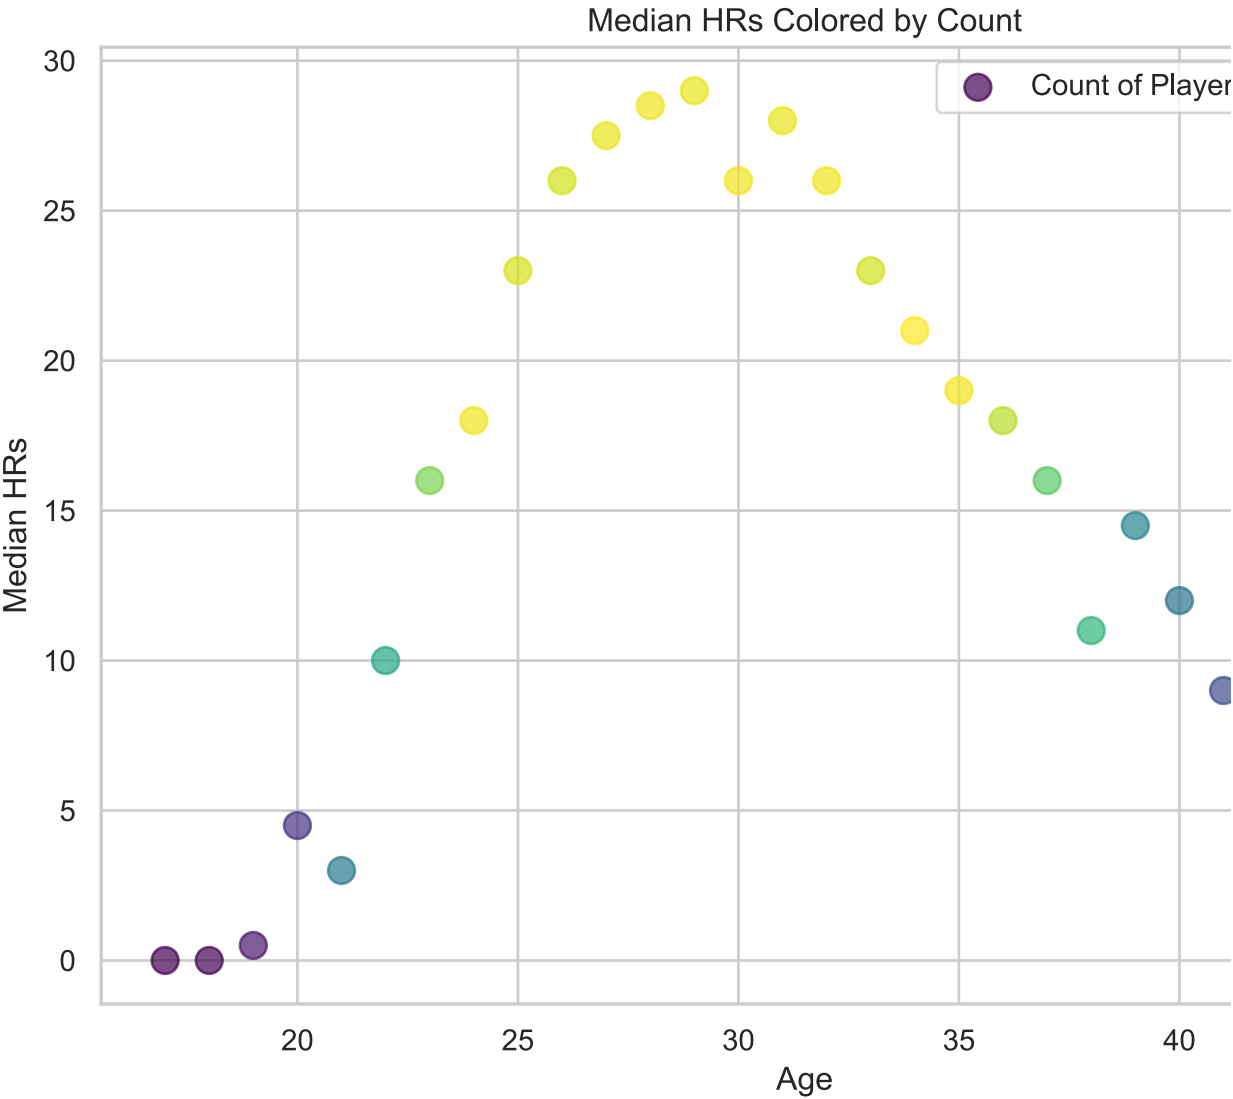

```
In [249]: 1 primeHRs = pd.merge(counts_df, medianHRs_df, on = "Age", how = "inner")
          2 primeHRs
```

Out[249]:

	Age	Count	MedianHRs
0	17	1	0.0
1	18	3	0.0
2	19	12	0.5
3	20	28	4.5
4	21	72	3.0
5	22	108	10.0
6	23	143	16.0
7	24	174	18.0
8	25	167	23.0
9	26	166	26.0
10	27	172	27.5
11	28	174	28.5
12	29	172	29.0
13	30	174	26.0
14	31	173	28.0
15	32	174	26.0
16	33	165	23.0
17	34	178	21.0
18	35	175	19.0
19	36	160	18.0
20	37	133	16.0
21	38	117	11.0
22	39	82	14.5
23	40	69	12.0
24	41	43	9.0
25	42	29	4.0
26	43	10	9.0
27	44	6	3.5
28	45	2	2.5
29	46	1	1.0

```
In [250]: 1 # Set up the figure and axes for plotting
2 plt.figure(figsize=(10, 6))
3
4 # Create a scatter plot for median HRs with varying colors based on cou
5 scatter = plt.scatter(
6     x=primeHRs['Age'],
7     y=primeHRs['MedianHRs'],
8     c=primeHRs['Count'], # Use count for coloring
9     cmap='viridis', # Choose a colormap
10    s=100, # Marker size
11    alpha=0.7, # Transparency
12    label='Count of Players in the Calculation'
13 )
14
15 # Add colorbar
16 cbar = plt.colorbar(scatter)
17 cbar.set_label('Count')
18
19 # Add labels and title
20 plt.xlabel('Age')
21 plt.ylabel('Median HRs')
22 plt.title('Median HRs Colored by Count')
23 plt.legend()
24
25 # Show the plot
26 plt.tight_layout()
27 plt.show()
```

```
/var/folders/cj/jcn0rh52g397xd7gy7ndwnc0000gn/T/ipykernel_51952/14756100
78.py:16: MatplotlibDeprecationWarning: Auto-removal of grids by pcolor()
and pcolormesh() is deprecated since 3.5 and will be removed two minor re
leases later; please call grid(False) first.
    cbar = plt.colorbar(scatter)
```



Linear Regression on Team Winning

```
In [251]: defining predictors for linear regression  
the idea is that I scale stats based on games played  
query = """  
SELECT W * (G/162) AS "Wins", HR * (G/162) AS "HRs", BB * (G/162) AS "BBs",  
FROM team_stats  
WHERE G >= 155;  
7  
8  
9  
pd.read_sql(query, engine)  
10
```

Out[251]:

	Wins	HRs	BBs	Es	BA	SBs	ERA
0	84.0000	161.0000	617.0000	126.0000	272.0327	126.0000	4.52
1	85.0000	147.0000	510.0000	106.0000	271.7584	93.0000	4.49
2	70.0000	158.0000	511.0000	106.0000	255.5515	71.0000	4.79
3	82.0000	236.0000	608.0000	134.0000	279.6731	93.0000	5.00
4	75.0000	158.0000	494.0000	103.0000	260.6738	116.0000	4.20
...
1733	86.0000	148.0000	630.0000	140.0000	250.5967	52.0000	3.49
1734	70.0000	138.0000	635.0000	116.0000	238.4615	72.0000	3.80
1735	61.8333	84.4074	564.3519	138.3889	230.4348	66.7407	3.70
1736	48.7963	34.4444	354.0123	422.9012	270.6867	188.4877	4.52
1737	51.6667	44.9691	334.8765	385.5864	271.8798	168.3951	4.93

1738 rows x 7 columns

```
In [252]: 1 #Linear regression with 4 features
2 data = df[['BBs', 'Es', 'BA', 'ERA']]
3 x = data.to_numpy() # convert to numpy array
4 X = sm.add_constant(x) # add a column of all 1s
5 y = df[["Wins"]].to_numpy()
6 model = sm.OLS(y,X) #run OLS
7 results = model.fit()
8 print(results.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          y      R-squared:
0.741
Model:                OLS      Adj. R-squared:
0.740
Method:              Least Squares      F-statistic:
1237.
Date:                Wed, 23 Aug 2023      Prob (F-statistic):
0.00
Time:                20:05:10      Log-Likelihood:          -5
645.3
No. Observations:      1738      AIC:                1.13
0e+04
Df Residuals:          1733      BIC:                1.13
3e+04
Df Model:              4
Covariance Type:      nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	
0.975]						

const	10.8395	2.928	3.702	0.000	5.097	1
6.582						
x1	0.0501	0.002	24.629	0.000	0.046	
0.054						
x2	-0.1097	0.004	-30.430	0.000	-0.117	-
0.103						
x3	0.4414	0.011	38.548	0.000	0.419	
0.464						
x4	-14.5919	0.258	-56.631	0.000	-15.097	-1
4.086						

```

=====
=====
Omnibus:              47.057      Durbin-Watson:
1.343
Prob(Omnibus):        0.000      Jarque-Bera (JB):          8
8.628
Skew:                 0.183      Prob(JB):                5.6
8e-20
Kurtosis:             4.044      Cond. No.                1.1
7e+04
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.17e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [253]: 1 #Linear regression with 6 features
2 data = df[['HRs', 'SBs', 'ERA']]
3 x = data.to_numpy() # convert to numpy array
4 X = sm.add_constant(x) # add a column of all 1s
5 y = df[["Wins"]].to_numpy()
6 model = sm.OLS(y,X) #run OLS
7 results = model.fit()
8 print(results.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          y      R-squared:
0.576
Model:                OLS      Adj. R-squared:
0.575
Method:              Least Squares      F-statistic:
785.6
Date:                Wed, 23 Aug 2023      Prob (F-statistic):      1.63
e-322
Time:                20:05:10      Log-Likelihood:      -6
072.2
No. Observations:      1738      AIC:      1.21
5e+04
Df Residuals:          1734      BIC:      1.21
7e+04
Df Model:              3
Covariance Type:      nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	
0.975]						

const	108.0434	1.355	79.751	0.000	105.386	11
0.701						
x1	0.1686	0.004	39.649	0.000	0.160	
0.177						
x2	0.0371	0.004	8.912	0.000	0.029	
0.045						
x3	-13.7936	0.338	-40.848	0.000	-14.456	-1
3.131						

```

=====
=====
Omnibus:              0.746      Durbin-Watson:
1.216
Prob(Omnibus):        0.689      Jarque-Bera (JB):
0.704
Skew:                 -0.048      Prob(JB):
0.703
Kurtosis:             3.021      Cond. No.      1.2
2e+03
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.22e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Autoregression on Wins

```
In [254]: 1 #Scale wins for prev year and curr year to 162
2 query = """
3     SELECT t1.W * (t1.G / 162) AS "PrevWins", t2.W * (t2.G / 162) AS "C
4     FROM team_stats AS t1
5     INNER JOIN team_stats AS t2
6     ON t1.teamID = t2.teamID AND t1.year = t2.year - 1
7     WHERE t1.G >= 155 AND t2.G >= 155;
8     """
9
10 df = pd.read_sql(query, engine)
11 df
```

Out[254]:

	PrevWins	CurrWins
0	84.0000	85.0000
1	85.0000	70.0000
2	70.0000	82.0000
3	82.0000	75.0000
4	75.0000	99.0000
...
1434	75.5309	64.5988
1435	64.5988	86.0000
1436	86.0000	70.0000
1437	70.0000	61.8333
1438	48.7963	51.6667

1439 rows × 2 columns

```
In [255]: 1 #Run linear regression
2 x = df["PrevWins"].to_numpy()
3 y = df["CurrWins"].to_numpy()
4 X = sm.add_constant(x) # add a column of all 1s
5 model = sm.OLS(y,X) #run OLS
6 results = model.fit()
7 print(results.summary())
8 bhat, ahat = results.params #Grab values
9 sigma_eps_hat = np.sqrt(results.mse_resid)
10
11 # Graph it
12 fig,ax = plt.subplots(figsize=(6,3))
13 ax.plot(x,ahat*x + bhat, '-', label="Model Trendline") # graph line of
14 ax.plot(x,y,"ko", label="Actual Wins")
15 ax.set_xlabel("Wins Year 1")
16 ax.set_ylabel("Wins Year 2")
17 ax.legend(loc="upper left", bbox_to_anchor=(1, 1))
18
19 #Graph residuals
20 r = results.resid
21 y_mean_pred = y - r
22 fig,ax = plt.subplots(figsize=(5,2))
23 ax.plot(y_mean_pred,r,"o", alpha = 0.25) # Plot the residuals
24 ax.plot(y_mean_pred,np.zeros(len(y)),"-") # Plot the reference line of
25 ax.set_xlabel("Wins")
26 ax.set_ylabel("Residual")
```

OLS Regression Results

```

=====
=====
Dep. Variable:          y    R-squared:
0.293
Model:                OLS    Adj. R-squared:
0.292
Method:              Least Squares    F-statistic:
594.5
Date:                Wed, 23 Aug 2023    Prob (F-statistic):        3.60
e-110
Time:                20:05:10    Log-Likelihood:            -5
346.5
No. Observations:    1439    AIC:                            1.07
0e+04
Df Residuals:        1437    BIC:                            1.07
1e+04
Df Model:              1
Covariance Type:      nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	
0.975]						

const	37.3423	1.788	20.882	0.000	33.834	4
0.850						
x1	0.5377	0.022	24.382	0.000	0.494	
0.581						

```

=====
=====
Omnibus:              6.742    Durbin-Watson:
2.130
Prob(Omnibus):        0.034    Jarque-Bera (JB):
5.738
Skew:                 -0.082    Prob(JB):
0.0567
Kurtosis:             2.737    Cond. No.
553.
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Out[255]: Text(0, 0.5, 'Residual')

